

# Assignment 2

## Kaggle Competition GTSRB

### Aishwarya Budhkar

#### 1) Summary

I tried to experiment with the following models

- Used 3- layer Convolutional Network as base test
- Used ResNet 18. Modified the same file for ResNet with STN
- Used VGG. Tried to implement VGG. Accuracy achieved is 97%
- Used Spatial Transformer Network modules with modified layers
- Used Spatial Transformer Network modules with ResNet18

#### 2) Introduction

Assignment is to develop and train models which will be used to classify signs in German Traffic Signs Recognition Benchmark [1] dataset. Starter code and test and train images are provided.

#### 3) Preprocessing

- Rescale- Rescaled the images to 48x48 and 64x64. When rescaled the images to 48X48, the results were better than 32x32.
- Transform- Tried to use rotate Transform for ResNet18 but didn't get good results.

#### 4) Models Implemented

##### Model 1:

Tried the basic 2-layer model given on course website. Trained on data to achieve accuracy 92%. Then trained a simple 3-layer convolutional network with 2 fully connected layers and high number of channels (100, 150 and 250). Achieved a test accuracy of around 97.7% with this model.

Optimizer- Stochastic Gradient Descent

Learning rate – 0.1

##### Improvement:


Added batch normalization layers to this CNN. Changed learning rate to 0.01.

Increased Image size to 48. Got best accuracy of 98.828%


Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 100, 42, 42]	14,800
BatchNorm2d-2	[-1, 100, 42, 42]	200
Conv2d-3	[-1, 150, 18, 18]	240,150
BatchNorm2d-4	[-1, 150, 18, 18]	300
Conv2d-5	[-1, 250, 6, 6]	600,250
BatchNorm2d-6	[-1, 250, 6, 6]	500
Linear-7	[-1, 300]	675,300
Linear-8	[-1, 43]	12,943

Total params: 1,544,443  
 Trainable params: 1,544,443  
 Non-trainable params: 0


Training accuracy:




Validation accuracy:



Training loss:



Validation loss:



## Model 2:


Increased number of neurons in output layer. Trained a simple 3-layer convolutional network with 2 fully connected layers and high number of channels (150, 200 and 300). Increased Image size to 64. I got an accuracy of 97.672%.

Optimizer - SGD


Learning rate- 0.01

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 100, 58, 58]	14,800
BatchNorm2d-2	[-1, 100, 58, 58]	200
Conv2d-3	[-1, 150, 26, 26]	240,150
BatchNorm2d-4	[-1, 150, 26, 26]	300
Conv2d-5	[-1, 250, 10, 10]	600,250
BatchNorm2d-6	[-1, 250, 10, 10]	500
Linear-7	[-1, 350]	2,187,850
Linear-8	[-1, 43]	15,093
Total params: 3,059,143		
Trainable params: 3,059,143		
Non-trainable params: 0		


Training accuracy:




Validation accuracy:



Training loss:



Validation loss:



Increase in the image size didn't help to increase the accuracy. Training was slower. Tried to change the momentum to 0.7. But the accuracy didn't change much though the convergence was better.

### Model 3:

Spatial Transformer Network – Currently, the state of the art results are held by a model [2] which uses Spatial Transformer Networks [3]. Both local and global features through spatial normalization (both locally and globally). STNs are used to avoid data augmentation and take care of color and spatial variance of image. Performs explicit geometric transformation for images.

Added batch normalization and drop out to improve the accuracy by normalizing the inputs.

Learning rate 0.001

Learning rate decay: 0.00001

Optimizer- Adam


Accuracy: 98.986%

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 10, 52, 52]	40
Dropout2d-2	[-1, 10, 52, 52]	0
BatchNorm2d-3	[-1, 10, 52, 52]	20
Conv2d-4	[-1, 3, 56, 56]	33
Dropout2d-5	[-1, 3, 56, 56]	0
BatchNorm2d-6	[-1, 3, 56, 56]	6
Conv2d-7	[-1, 16, 52, 52]	2,368
Conv2d-8	[-1, 32, 24, 24]	12,832
Conv2d-9	[-1, 64, 12, 12]	18,496
Linear-10	[-1, 128]	295,040
Linear-11	[-1, 64]	8,256
Linear-12	[-1, 6]	390
Conv2d-13	[-1, 16, 48, 48]	1,216


Dropout2d-14	[-1, 16, 48, 48]	0
BatchNorm2d-15	[-1, 16, 48, 48]	32
Conv2d-16	[-1, 32, 48, 48]	12,832
Dropout2d-17	[-1, 32, 48, 48]	0
BatchNorm2d-18	[-1, 32, 48, 48]	64
Conv2d-19	[-1, 64, 48, 48]	51,264
Dropout2d-20	[-1, 64, 48, 48]	0
BatchNorm2d-21	[-1, 64, 48, 48]	128
Conv2d-22	[-1, 96, 48, 48]	153,696
Dropout2d-23	[-1, 96, 48, 48]	0
BatchNorm2d-24	[-1, 96, 48, 48]	192
Conv2d-25	[-1, 128, 48, 48]	307,328
Dropout2d-26	[-1, 128, 48, 48]	0
BatchNorm2d-27	[-1, 128, 48, 48]	256
Conv2d-28	[-1, 192, 48, 48]	614,592
Dropout2d-29	[-1, 192, 48, 48]	0
BatchNorm2d-30	[-1, 192, 48, 48]	384
Conv2d-31	[-1, 256, 24, 24]	1,229,056
Dropout2d-32	[-1, 256, 24, 24]	0
BatchNorm2d-33	[-1, 256, 24, 24]	512
Conv2d-34	[-1, 128, 12, 12]	819,328
Dropout2d-35	[-1, 128, 12, 12]	0
BatchNorm2d-36	[-1, 128, 12, 12]	256
Conv2d-37	[-1, 64, 12, 12]	204,864
Dropout2d-38	[-1, 64, 12, 12]	0
BatchNorm2d-39	[-1, 64, 12, 12]	128
Linear-40	[-1, 43]	2,795
Dropout2d-41	[-1, 43]	0

=====  
Total params: 3,736,404  
Trainable params: 3,736,404  
Non-trainable params: 0


Training accuracy:




Testing accuracy:



Training loss:



Testing loss:



#### Model 4:

Tried to implement ResNet 18.

Learning rate – 0.0001

Optimizer-Adam

Accuracy achieved was 98.957%.

So, tried to use Spatial Transform Network with Resnet 18:

Image size: 48

Got Test accuracy: 98.337%

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 10, 52, 52]	40
Dropout2d-2	[-1, 10, 52, 52]	0
BatchNorm2d-3	[-1, 10, 52, 52]	20
Conv2d-4	[-1, 3, 56, 56]	33
Dropout2d-5	[-1, 3, 56, 56]	0
BatchNorm2d-6	[-1, 3, 56, 56]	6
Conv2d-7	[-1, 16, 52, 52]	2,368
Conv2d-8	[-1, 32, 24, 24]	12,832
Conv2d-9	[-1, 64, 12, 12]	18,496
Linear-10	[-1, 128]	295,040
Linear-11	[-1, 64]	8,256
Linear-12	[-1, 6]	390
Conv2d-13	[-1, 64, 24, 24]	9,408
BatchNorm2d-14	[-1, 64, 24, 24]	128

MaxPool2d-15	[-1, 64, 12, 12]	0
Conv2d-16	[-1, 64, 12, 12]	36,864
BatchNorm2d-17	[-1, 64, 12, 12]	128
Conv2d-18	[-1, 64, 12, 12]	36,864
BatchNorm2d-19	[-1, 64, 12, 12]	128
BasicBlock-20	[-1, 64, 12, 12]	0
Conv2d-21	[-1, 64, 12, 12]	36,864
BatchNorm2d-22	[-1, 64, 12, 12]	128
Conv2d-23	[-1, 64, 12, 12]	36,864
BatchNorm2d-24	[-1, 64, 12, 12]	128
BasicBlock-25	[-1, 64, 12, 12]	0
Conv2d-26	[-1, 128, 6, 6]	73,728
BatchNorm2d-27	[-1, 128, 6, 6]	256
Conv2d-28	[-1, 128, 6, 6]	147,456
BatchNorm2d-29	[-1, 128, 6, 6]	256
Conv2d-30	[-1, 128, 6, 6]	8,192
BatchNorm2d-31	[-1, 128, 6, 6]	256
BasicBlock-32	[-1, 128, 6, 6]	0
Conv2d-33	[-1, 128, 6, 6]	147,456
BatchNorm2d-34	[-1, 128, 6, 6]	256
Conv2d-35	[-1, 128, 6, 6]	147,456
BatchNorm2d-36	[-1, 128, 6, 6]	256
BasicBlock-37	[-1, 128, 6, 6]	0
Conv2d-38	[-1, 256, 3, 3]	294,912
BatchNorm2d-39	[-1, 256, 3, 3]	512
Conv2d-40	[-1, 256, 3, 3]	589,824
BatchNorm2d-41	[-1, 256, 3, 3]	512
Conv2d-42	[-1, 256, 3, 3]	32,768
BatchNorm2d-43	[-1, 256, 3, 3]	512
BasicBlock-44	[-1, 256, 3, 3]	0
Conv2d-45	[-1, 256, 3, 3]	589,824
BatchNorm2d-46	[-1, 256, 3, 3]	512
Conv2d-47	[-1, 256, 3, 3]	589,824
BatchNorm2d-48	[-1, 256, 3, 3]	512
BasicBlock-49	[-1, 256, 3, 3]	0
Conv2d-50	[-1, 512, 2, 2]	1,179,648
BatchNorm2d-51	[-1, 512, 2, 2]	1,024
Conv2d-52	[-1, 512, 2, 2]	2,359,296
BatchNorm2d-53	[-1, 512, 2, 2]	1,024
Conv2d-54	[-1, 512, 2, 2]	131,072
BatchNorm2d-55	[-1, 512, 2, 2]	1,024
BasicBlock-56	[-1, 512, 2, 2]	0
Conv2d-57	[-1, 512, 2, 2]	2,359,296
BatchNorm2d-58	[-1, 512, 2, 2]	1,024
Conv2d-59	[-1, 512, 2, 2]	2,359,296
BatchNorm2d-60	[-1, 512, 2, 2]	1,024
BasicBlock-61	[-1, 512, 2, 2]	0
AvgPool2d-62	[-1, 512, 1, 1]	0
Linear-63	[-1, 43]	22,059

=====


Total params: 11,536,052

Trainable params: 11,536,052


Non-trainable params: 0

-----


Training accuracy:




Validation accuracy:



Training loss:



Validation loss:



The model didn't perform as well on test data.

Tried to decrease the learning rate to 5e-5. But the accuracy didn't improve.

## 5) Results

Model	Image size	Epochs	Learning rate	Optimizer	Accuracy
2-layer CNN	32x32	20	0.1	SGD	92
3-layer CNN	32x32	20	0.01	SGD	94.568
3-layer CNN with Batch Normalization	48x48	100	0.01	SGD	98.828
3-layer CNN with Batch Normalization	48x48	150	0.01	SGD	98.828
3-layer CNN with Batch Normalization	64x64	100	0.01	SGD	97.672



CNN with Spatial Transformer Network	48x48	100	0.01	SGD	98.749
CNN with Spatial Transformer Network	48x48	100	0.001	SGD	98.875
CNN with Spatial Transformer Network	48x48	100	0.001	Adam	98.986
CNN with Spatial Transformer Network	48x48	150	0.00005	Adam	98.337
ResNet with Spatial Transform Network	48x48	100	0.001	Adam	98.606
ResNet with Spatial Transform Network	48x48	150	0.0001	Adam	98.083
ResNet with Spatial Transform Network	48x48	150	0.00005	Adam	98.004
ResNet	48x48	100	0.01	SGD	97.957

## 6) Conclusion

The images in dataset have severe flaws such as blur, shadow, obstructions, bad lightening. There are distortions like rotation, translation, etc. These are difficult for human to recognize. Spatial Transformer networks are good for classification tasks with lot of distortion in data. As, STNs take care of color and spatial variance and avoid augmentation, with learned color and spatial variance the model can adapt to image's characteristics.

### Improvement:

More experimentation with hyper parameters can lead to better results. Also, multiple STNs [4] can be added to get better accuracy. As the current model is trained on very limited dataset, it is unable to recognize other signs which it didn't encounter. With more data, the accuracy might improve.

## 7) References

- [1] Johannes Stallkamp, Marc Schlipsing, Jan Salmen, and Christian Igel. The german traffic sign recognition benchmark: a multi-class classification competition. In Neural Networks (IJCNN), The 2011 International Joint Conference on, pages 1453–1460. IEEE, 2011.
- [2] Mrinal Haloi. Traffic sign classification using deep inception based convolutional networks. CoRR, abs/1511.02992, 2015.
- [3] Max Jaderberg, Karen Simonyan, Andrew Zisserman, and Koray Kavukcuoglu. Spatial transformer networks. CoRR, abs/1506.02025, 2015.
- [4] A. Arcos-Garcia, J. Alvarez-Garcia, and Luis M.Soria-Morillo. Deep neural network for traffic sign recognition systems: An analysis of spatial transformers and stochastic optimisation methods. In Neural Networks, pages 158–165, 2018.